

Kdtree: Design and Implementation

Russ Hamilton

May 26, 2016

Overview

- Provides classes and functions for
 - Reading k-dimensional points from a CSV from a stream
 - Generating a kd-tree from the set of points
 - kd-tree is optimized for nearest neighbor search
 - Serializing the kd-tree to a stream
 - Deserializing the kd-tree from a stream
 - Efficiently find the nearest neighbor in the tree to an arbitrary point

Tree Representation

- Point coordinates
 - Kept in flat array
 - Calculations refer to points indirectly by using their index in the array
 - Reduces the memory footprint and amount of copying required to build the tree
- Tree structure
 - Array of nodes (with the same length as the array of points)
 - Each node keeps track of
 - The axis to split over
 - The index of this node
 - The index of its left and right children
 - Keeping indices instead of references make serialization easy

Building the Tree

- Recursive algorithm
 - Takes the list of points
 - Finds the axis with largest difference
 - Sorts points (indirectly) along this axis
 - Splits the list at the median point
 - Creates a node
 - Left child is the subtree generated by the list of points before the median
 - Right child is the subtree generated by the list of points after the median

Building the Tree

- Complexity - $O(kn \log n + n \log^2 n)$
 - Finding the axis with the largest difference $O(kn)$
 - Assume `std::minmax_element` is $O(n)$
 - Sort points along a single axis $O(n \log n)$
 - Recursive algorithm runs $O(\log n)$ times
- Potential optimization (not implemented)
 - Pre-sort the list of points across all dimensions
 - Makes finding axis with largest difference $O(k)$
 - No sort required
 - Overall complexity $O(kn \log n)$

Finding the Nearest Neighbor

- Recursive algorithm to find nearest neighbor within bound
 - If current node is closest so far, update the closest
 - Recurse with subtree that would contain the target (if any), update the closest
 - Recurse with other subtrees if the region could contain a point closer than the current closest, update the closest

Finding the Nearest Neighbor

- Complexity – $O(k^2 \log n)$
 - Point distance and region calculations – $O(k)$
 - Number of subtrees containing the target – $O(\log n)$
 - Number of subtrees not containing the target that contain points closer to the target than the subtree containing the target
 - Worst case $O(n)$
 - Typical case $O(k \log n)$

Summary

- Provides moderately efficient tree generation
- Provides efficient exact nearest neighbor search